

# Computer Graphics

---

**Prof. Feng Liu**

**Fall 2016**

<http://www.cs.pdx.edu/~fliu/courses/cs447/>

**10/03/2016**

# Announcements

---

- Free Textbook: Linear Algebra
  - By Jim Hefferon
  - <http://joshua.smcvt.edu/linalg.html/>
- Homework 1 due in class on Oct. 05
- Project 1 is available on course website
  - due 5pm October 28

# Last Time

---

## □ Color

- The principle of trichromacy says that any spectrum can be matched using three primaries (but sometimes you have to subtract a primary)
- A color system consist of primaries and color matching functions that are used to determine how much of each primary is needed to match a spectrum
- RGB, CIE XYZ, HSV are some examples of color systems
- Linear color spaces make it easy to convert between colors - matrix multiply
- Color calibration is an important step to achieving accurate color


## □ Today

- Perceptually linear (uniform) color spaces make distances between colors meaningful

# RGB Color Space

---

## □ Demo




# HSV Color Space (Alvy Ray Smith, 1978)

---

- Hue: the color family: red, yellow, blue...
  - Saturation: The purity of a color: white is totally unsaturated
  - Value: The intensity of a color: white is intense, black isn't
  - Space looks like a cone
    - Parts of the cone can be mapped to RGB space
  - Not a linear space, so no linear transform to take RGB to HSV
    - But there is an algorithmic transform
-

# HSV Color Space

---



# Linear Space vs. Perceptually Uniform

---


## □ Linear Space: RGB, CIE XYZ

- The principle of trichromacy means that the colors displayable are all the linear combination of primaries
  - HSV is not a linear space
- Matrix multiplication
- Easy to convert between colors
- Not perceptually linear

## □ Perceptually Uniform space

- Computational consuming
- Make color distance meaningful
- CIE  $u'v'$ : a good approximation


# MacAdam Ellipses



- Refer to the region which contains all colors which are indistinguishable
- Scaled by a factor of 10 and shown on CIE xy color space
- If you are shown two colors, one at the center of the ellipse and the other inside it, you cannot tell them apart
- Only a few ellipses are shown, but one can be defined for every point



# CIE u'v' Space



- CIE u'v' is a non-linear color space where color differences are more uniform
- Note that now ellipses look more **like** circles
- The third coordinate is the original Z from XYZ

$$\begin{bmatrix} u' \\ v' \end{bmatrix} = \frac{1}{X + 15Y + 3Z} \begin{bmatrix} 4X \\ 9Y \end{bmatrix}$$

# Today

---

- ☐ Ink
- ☐ Image file formats
- ☐ Color quantization
- ☐ Programming tutorial 2
  - How to use FLTK within Visual Studio

# Ink

---

- Ink is thought of as *adsorbing* particles
  - You see the color of the paper, filtered by the ink
  - Combining inks adsorbs more color, so subtractive color
    - White paper - red - blue = green
  - The color and texture of the paper affects the color of the image

# Subtractive mixing

---

- Common inks:  $\text{Cyan} = \text{White} - \text{Red}$ ;  
 $\text{Magenta} = \text{White} - \text{Green}$ ;  $\text{Yellow} = \text{White} - \text{Blue}$ 
  - cyan, magenta, yellow, are how the inks look when printed
- For good inks, matching is linear:
  - $\text{C} + \text{M} + \text{Y} = \text{White} - \text{White} = \text{Black}$
  - $\text{C} + \text{M} = \text{White} - \text{Red} - \text{Green} = \text{Blue}$
  - How to make a red mark?

# Subtractive mixing

---

- Common inks: Cyan=White-Red;  
Magenta=White-Green; Yellow=White-Blue
  - cyan, magenta, yellow, are how the inks look when printed
- For good inks, matching is linear:
  - $C+M+Y = \text{White} - \text{White} = \text{Black}$
  - $C+M = \text{White} - \text{Red} - \text{Green} = \text{Blue}$
  - How to make a red mark?
- Usually require CMY and Black, because colored inks are more expensive, and registration is hard
  - Registration is the problem of making drops of ink line up

# Calibrating a Printer

---

- If the inks (think of them as primaries) are linear, there exists a 3x3 matrix and an offset to take RGB to CMY
  - For example, if an RGB of (1,0,0) goes to CMY of (0,1,1); (0,1,0)→(1,0,1); and (0,0,1)→(1,1,0), then the matrix is

$$\begin{bmatrix} c \\ m \\ y \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} + \begin{bmatrix} -1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & -1 \end{bmatrix} \begin{bmatrix} r \\ g \\ b \end{bmatrix}$$

- To calibrate your printer, you find out exactly what the numbers in the matrix should be
  - Print with cyan ink only and match the color with RGB, repeat with magenta and yellow, use the results to determine the matrix

# Image File Formats

---


- How big is the image?
  - All files in some way store width and height
- How is the image data formatted?
  - Is it a black and white image, a grayscale image, a color image, an *indexed color* image?
  - How many bits per pixel?
- What other information?
  - Color tables, compression codebooks, creator information...
- All image formats are a trade-off between ease of use, size of file, and quality of reproduction

# The Simplest File

- ❑ Assumes that the color depth is known and agreed on
- ❑ Store width, height, and data for every pixel in sequence
- ❑ This is how you normally store an image in memory

0 <sub>r,g,b</sub>	1 <sub>r,g,b</sub>	2 <sub>r,g,b</sub>
3 <sub>r,g,b</sub>	4 <sub>r,g,b</sub>	5 <sub>r,g,b</sub>
6 <sub>r,g,b</sub>	7 <sub>r,g,b</sub>	8 <sub>r,g,b</sub>

```
class Image {  
    unsigned int width;  
    unsigned int height;  
    unsigned char *data;  
}
```



- ❑ Unsigned because width and height are positive, and unsigned char because it is the best type for raw 8 bit data
- ❑ Note that you require some implicit scheme for laying out a rectangular array into a linear one



# Indexed Color

---

- ❑ 24 bits per pixel (8-red, 8-green, 8-blue) are expensive to transmit and store
- ❑ It must be possible to represent all those colors, but *not in the same image*
- ❑ Solution: Indexed color
  - Assume  $k$  bits per pixel (typically 8)
  - Define a *color table* containing  $2^k$  colors (24 bits per color)
  - Store the *index* into the table for each pixel (so store  $k$  bits for each pixel, instead of 24 bits)
  - Once common in hardware, now an artifact (256 color displays)

# Indexed Color

---


Color Table

0	
1	
2	
3	
4	
5	
6	
7	

Pixel Data

4	3	0	2
1	7	4	5
3	7	6	5
2	2	1	1

Image



Only makes sense if you have lots of pixels and not many colors

# Image Compression

---

- ❑ Indexed color is one form of image compression
  - Special case of *vector quantization* - in color space, reducing the range of available colors
- ❑ Alternative 1: Store the image in a simple format and then compress with your favorite compressor
  - Doesn't exploit image specific information
  - Doesn't exploit perceptual shortcuts
- ❑ Two historically common compressed file formats: GIF and JPEG
  - GIF should now be replaced with PNG, because GIF is patented and the owner started enforcing the patent
    - ❑ Patent expired recently?

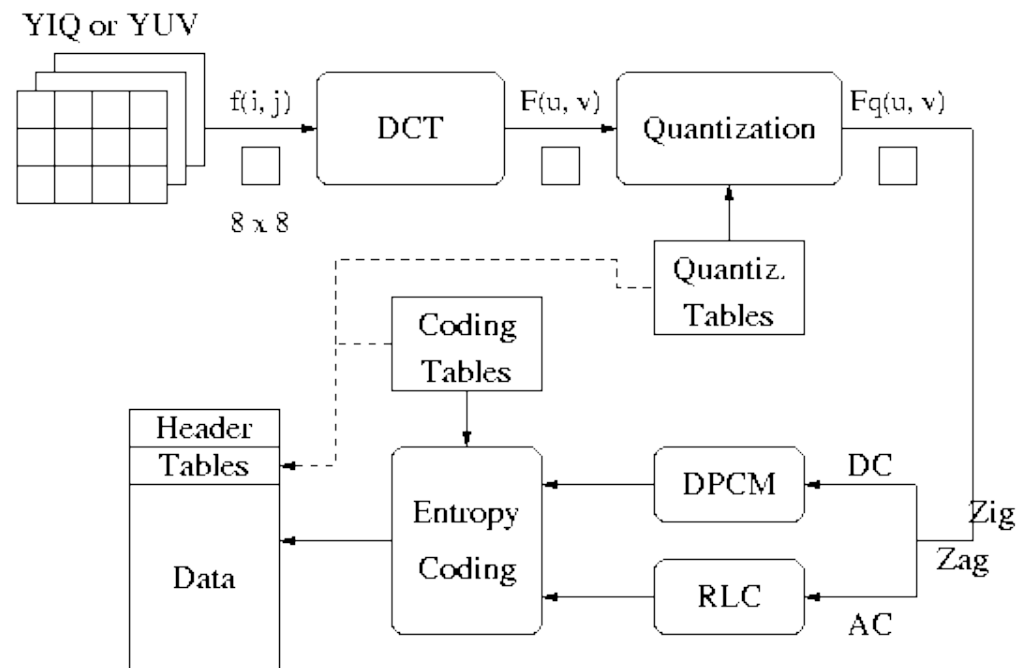
# GIF

---

- ❑ Header - Color Table - Image Data - Extensions
- ❑ Header gives basic information such as size of image and size of color table
- ❑ Color table gives the colors found in the image
  - Biggest it can be is 256 colors, smallest is 2
- ❑ Image data is LZW compressed color indices
- ❑ To create a GIF:
  - Choose colors
  - Create an array of color indices
  - Compress it with LZW

# JPEG


- ❑ Multi-stage process intended to get very high compression with controllable quality degradation
- ❑ Start with YIQ color



# Discrete Cosine Transform

---

- A transformation to convert from the *spatial* to *frequency* domain - done on 8x8 blocks
- Why? Humans have varying sensitivity to different frequencies, so it is safe to throw some of them away
- Basis functions:



# Quantization

---

- ❑ Reduce the number of bits used to store each coefficient by dividing by a given value
  - If you have an 8 bit number (0-255) and divide it by 8, you get a number between 0-31 (5 bits = 8 bits - 3 bits)
  - Different coefficients are divided by different amounts
  - Perceptual issues come in here
- ❑ Achieves the greatest compression, but also quality loss
- ❑ “Quality” knob controls how much quantization is done

# Entropy Coding

---

- Standard lossless compression on quantized coefficients
  - Delta encode the DC components
  - Run length encode the AC components
    - Lots of zeros, so store number of zeros then next value
  - Huffman code the encodings



# Lossless JPEG With Prediction

---

- ❑ Predict what the value of the pixel will be based on neighbors
- ❑ Record error from prediction
  - Mostly error will be near zero
- ❑ Huffman encode the error stream
- ❑ Variation works really well for fax messages

# Today

---

- Ink
- Image file formats
- Color quantization
- Programming tutorial 2
  - How to use FLTK within Visual Studio

# Color Quantization

---

- The problem of reducing the number of colors in an image with minimal impact on appearance
  - Extreme case: 24 bit color to black and white
  - Less extreme: 24 bit color to 256 colors, or 256 grays
- Sub problems:
  - Decide which colors to use in the output (if there is a choice)
  - Decide which of those colors should be used for each input pixel

# Example (24 bit color)

---



# Uniform Quantization


---

- ❑ Break the color space into uniform cells
- ❑ Find the cell that each color is in, and map it to the center
- ❑ Equivalent to dividing each color by some number and taking the integer part
  - Say your original image is 24 bits color (8 red, 8 green, 8 blue)
  - Say you have 256 colors available, and you choose to use 8 reds, 8 greens and 4 blues ( $8 \times 8 \times 4 = 256$ )
  - Divide original red by 32, green by 32, and blue by 64
  - Some annoying details
- ❑ Generally does poorly because it fails to capture the distribution of colors
  - Some cells may be empty, and are wasted

# Uniform Quantization

---

- ❑ 8 bits per pixel in this image
- ❑ Note that it does very poorly on smooth gradients
- ❑ Normally the hardest part to get right, because lots of similar colors appear very close together
- ❑ Does this scheme use information from the image?



# Populosity Algorithm


---

- Build a color histogram: count the number of times each color appears
- Choose the  $n$  most commonly occurring colors
  - Typically group colors into *small* cells first using uniform quantization
- Map other colors to the closest chosen color
- Problem?

# Populosity Algorithm

---

- 8 bit image, so the most popular 256 colors






# Populosity Algorithm

---

- ❑ 8 bit image, so the most popular 256 colors
- ❑ Note that blue wasn't very popular, so the crystal ball is now the same color as the floor
- ❑ Populosity ignores rare but important colors!




# Median Cut (Clustering)

---

- View the problem as a *clustering* problem
  - Find groups of colors that are similar (a cluster)
  - Replace each input color with one representative of its cluster
- Many algorithms for clustering
- *Median Cut* is one: recursively
  - Find the “longest” dimension (r, g, b are dimensions)
  - Choose the median of the long dimension as a color to use
  - Split into two sub-clusters along the median plane, and recurse on both halves
- Works very well in practice


# Median Cut (Clustering)

---




# Median Cut (Clustering)

---




# Median Cut (Clustering)

---



# Median Cut (Clustering)


---



# Median Cut

---

- 8 bit image, so 256 colors
- Now we get the blue
- Median cut works so well because it divides up the color space in the “most useful” way



# Optimization Algorithms

---

- The quantization problem can be phrased as optimization
  - Find the set of colors and map that result in the lowest quantization error
- Several methods to solve the problem, but of limited use unless the number of colors to be chosen is small
  - It's expensive to compute the optimum
  - It's also a poorly behaved optimization



# Perceptual Problems


---

- While a good quantization may get close colors, humans still perceive the quantization
- Biggest problem: *Mach bands*
  - The difference between two colors is more pronounced when they are side by side and the boundary is smooth
  - This emphasizes boundaries between colors, even if the color difference is small
  - Rough boundaries are “averaged” by our vision system to give smooth variation

# Mach Bands in Reality

---

The floor appears banded



# Mach Bands in Reality

---

Still some banding even in this 24 bit image (the floor in the background)



# Dithering (Digital Halftoning)

---

- Mach bands can be removed by adding noise along the boundary lines
- General perceptive principle: replaced structured errors with noisy ones and people complain less
- Old industry dating to the late 1800's
  - Methods for producing grayscale images in newspapers and books

# Next Time

---

- ☐ Dithering
- ☐ Sampling
- ☐ Signal Processing