

Computer Graphics

Prof. Feng Liu

Fall 2016

<http://www.cs.pdx.edu/~fliu/courses/cs447/>

10/10/2016

Last Time

- ☐ Dithering
- ☐ Signal Processing

Today

- ☐ Filtering
- ☐ Resampling
- ☐ Project 1 due October 28


Filters

- A *filter* is something that attenuates or enhances particular frequencies
- Easiest to visualize in the frequency domain, where filtering is defined as multiplication:


$$H(\omega) = F(\omega) \times G(\omega)$$

- Here, F is the spectrum of the function, G is the spectrum of the filter, and H is the filtered function. Multiplication is point-wise


Qualitative Filters



Low-Pass Filtered Image



High-Pass Filtered Image




Filtering in the Spatial Domain

- Filtering the spatial domain is achieved by *convolution*

$$h(x) = f \otimes g = \int_{-\infty}^{\infty} f(u)g(x-u)du$$

- Qualitatively: Slide the filter to each position, x , then sum up the function multiplied by the filter at that position

Convolution Example



Convolution Theorem

- Convolution in the spatial domain is the same as multiplication in the frequency domain
 - Take a function, f , and compute its Fourier transform, F
 - Take a filter, g , and compute its Fourier transform, G
 - Compute $H=F \times G$
 - Take the inverse Fourier transform of H , to get h
 - Then $h=f \otimes g$
- Multiplication in the spatial domain is the same as convolution in the frequency domain

Filtering Images

$$I_{output}[x][y] = \sum_{i=-k/2}^{k/2} \sum_{j=-k/2}^{k/2} I_{input}[x+i][y+j]M[i+k/2][j+k/2]$$

- Work in the discrete spatial domain
- Convert the filter into a matrix, the *filter mask*
- Move the matrix over each point in the image, multiply the entries by the pixels below, then sum

Handling Boundaries


$$I_{output}[x][y] = \sum_{i=-k/2}^{k/2} \sum_{j=-k/2}^{k/2} I_{input}[x+i][y+j] M[i+k/2][j+k/2]$$

- ❑ At (0,0) for instance, you might need pixel data for (-1,-1), which doesn't exist
- ❑ Option 1: Make the output image smaller - don't evaluate pixels you don't have all the input for
- ❑ Option 2: Replicate the edge pixels
 - Equivalent to: $posn = x + i$; if ($posn < 0$) $posn = 0$; and so on for other indices
- ❑ Option 3: Reflect image about edge
 - Equivalent to: $posn = x + i$; if ($posn < 0$) $posn = -posn$; and similar for others


Box Filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Box filters smooth by averaging neighbors
- In frequency domain, keeps low frequencies and attenuates (reduces) high frequencies, so clearly a low-pass filter



Spatial: Box




Frequency: sinc


Bartlett Filter

$$\frac{1}{81} \begin{bmatrix} 1 & 2 & 3 & 2 & 1 \\ 2 & 4 & 6 & 4 & 2 \\ 3 & 6 & 9 & 6 & 3 \\ 2 & 4 & 6 & 4 & 2 \\ 1 & 2 & 3 & 2 & 1 \end{bmatrix}$$

- Triangle shaped filter in spatial domain
- In frequency domain, product of two box filters, so attenuates high frequencies more than a box




Spatial: Triangle (Box \otimes Box)



Frequency: sinc²

Constructing Masks: 1D

- Sample the filter, then normalize
- eg 1D Bartlett



- Can go to edge of pixel or middle of next: results are slightly different

Constructing Masks: 2D

- Multiply 2 1D masks together using *outer product*

$$M = mm^T, \text{ or } M[i][j] = m[i]m[j]$$


- M is 2D mask, m is 1D mask

	0.2	0.6	0.2
0.2	0.04	0.12	0.04
0.6	0.12	0.36	0.12
0.2	0.04	0.12	0.04

Gaussian Filter

$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

- Attenuates high frequencies even further
- In 2d, rotationally symmetric, so fewer artifacts



Constructing Gaussian Mask

- Use the binomial coefficients
 - Central Limit Theorem (probability) says that with more samples, binomial converges to Gaussian


$$\frac{1}{4} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

$$\frac{1}{16} \begin{array}{|c|c|c|c|c|} \hline 1 & 4 & 6 & 4 & 1 \\ \hline \end{array}$$

$$\frac{1}{64} \begin{array}{|c|c|c|c|c|c|c|} \hline 1 & 6 & 15 & 20 & 15 & 6 & 1 \\ \hline \end{array}$$

Constructing Gaussian Mask

- Use the binomial coefficients
 - Central Limit Theorem (probability) says that with more samples, binomial converges to Gaussian




High-Pass Filters

- A high-pass filter can be obtained from a low-pass filter
 - If we subtract the smoothed image from the original, we must be subtracting out the low frequencies
 - What remains must contain only the high frequencies
- High-pass masks come from matrix subtraction:
- eg: 3x3 Bartlett

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \frac{1}{16} \begin{bmatrix} -1 & -2 & -1 \\ -2 & 12 & -2 \\ -1 & -2 & -1 \end{bmatrix}$$

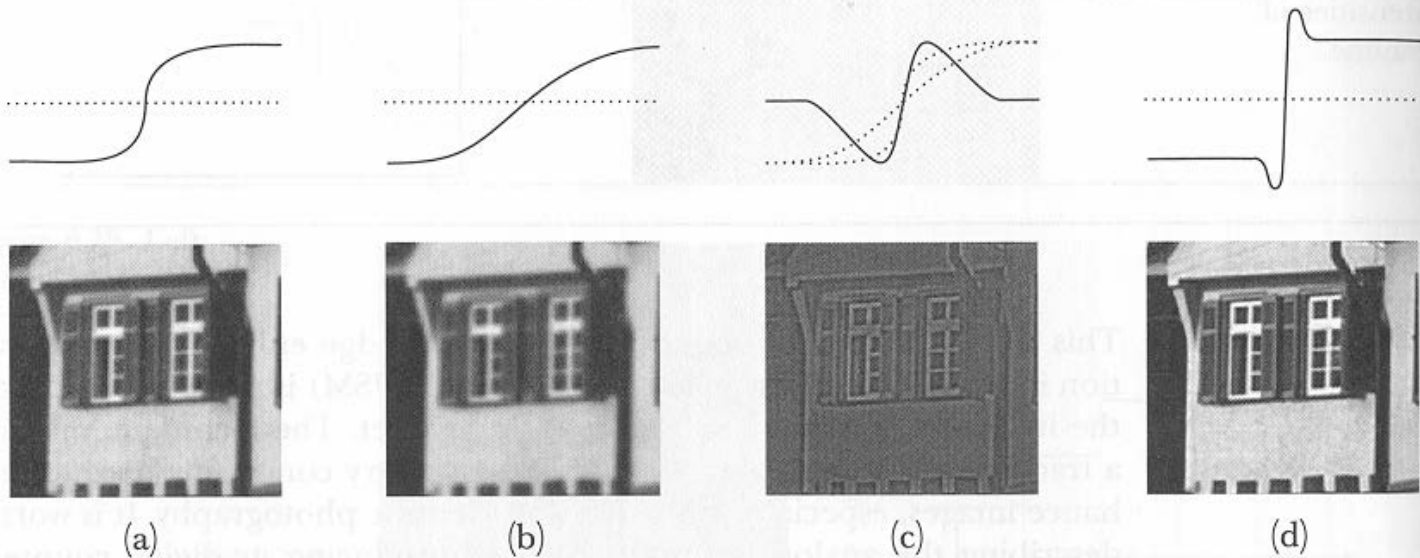
Edge Enhancement

- High-pass filters give high values at edges, low values in smooth regions
- Adding high frequencies back into the image enhances edges
- One approach:
 - $\text{Image} = \text{Image} + [\text{Image} - \text{smooth}(\text{Image})]$



Edge Enhancement

Figure 6.32. (a) Original image. (b) Blurred image. (c) Difference between first two. (d) Enhanced image.



Fixing Negative Values

- The negative values in high-pass filters can lead to negative image values
 - Most image formats don't support this
- Solutions:
 - Truncate: Chop off values below min or above max
 - Offset: Add a constant to move the min value to 0
 - Re-scale: Rescale the image values to fill the range (0,max)

Filtering and Color


- To filter a color image, simply filter each of R,G and B separately
- Re-scaling and truncating are more difficult to implement:
 - Adjusting each channel separately may change color significantly

Today

- ☐ Filtering
- ☐ Resampling
- ☐ Project 1 due October 28

Resampling


- ❑ Making an image larger is like sampling the original function at higher density
 - You need more pixels to represent the same thing, so a higher pixel density



- ❑ Reducing an image in size is like sampling at lower density
 - ❑ Generating new samples of the “same” function is called *resampling*
 - In theory, 2 steps: Reconstruction and sampling - but not in practice
 - ❑ Many other image manipulation tasks require resampling
-



General Scenario

- You are trying to create a new image of some form, and you need data from a particular place in the existing image
 - Always: Figure out where the new sample comes from in the original image




Resampling at a Point

- We want to reconstruct the original “function” at the required point
- We will use information from around the point to do this
- We do it using a filter
- Which filter?
 - We’ll look at Bartlett (triangular)
 - Other filters also work



Resampling at a Point

- Place a Bartlett filter at the required point
- Multiply the value of the neighbor by the filter at that point, and add them
 - Convolution with discrete samples
- The filter size is a parameter




- Say the filter is size 3, and you need the value at $x=5.75$
 - You need the image samples, $I(x)$, from $x=5$, $x=6$ and $x=7$
 - You need the filter value, $H(s)$, at $s=-0.75$, $s=0.25$ and $s=1.25$
 - Compute: $I(5)*H(-0.75)+I(6)*H(0.25)+I(7)*H(1.25)$
-

Functional Form for Filters

- Consider the Bartlett in 1D:

$$H_w(s) = \frac{2}{w} \left(1 - \frac{2|s|}{w} \right)$$



- To apply it at a point x_c and find the contribution from point x where the image has value $I(x)$

$$f(x) = \frac{2}{w} \left(1 - \frac{2|x - x_c|}{w} \right) I(x)$$

- Extends naturally to 2D:

$$f(x, y) = \frac{4}{w^2} \left(1 - \frac{2|x - x_c|}{w} \right) \left(1 - \frac{2|y - y_c|}{w} \right) I(x, y)$$

Common Operations

- Image scaling by a factor k (e.g. 0.5 = half size):

- To get x_{orig} given x_{new} divide by k :

$$(x_{orig}, y_{orig}) = \left(\frac{x_{new}}{k}, \frac{y_{new}}{k} \right)$$

- Image rotation by an angle θ :

$$x_{orig} = \cos(-\theta)x_{new} - \sin(-\theta)y_{new}$$

$$y_{orig} = \sin(-\theta)x_{new} + \cos(-\theta)y_{new}$$


- This rotates around the bottom left corner. It's up to you to figure out how to rotate about the center
- Be careful of radians vs. degrees: all C++ standard math functions take radians, but OpenGL functions take degrees

Ideal Image Resize

- To do ideal image resampling, we would reconstruct the original function based on the samples
- A *requirement* for *perfect* enlargement or size reduction
 - Almost never possible in practice, and we'll see why

An Reconstruction Example

- Say you have a sine function of a particular frequency
- And you sample it too sparsely
- You could draw a different sine curve through the samples



Some Intuition


- To reconstruct a function, you need to reconstruct every frequency component that's in it
 - This is in the frequency domain, but that's because it's easy to talk about “components” of the function
- But we've just seen that to accurately reconstruct high frequencies, you need more samples
- The effect on the previous slide is called aliasing
 - The correct frequency is *aliased* by the longer wavelength curve

Nyquist Frequency

- Aliasing cannot happen if you sample at a frequency that is twice the original frequency - the *Nyquist* sampling limit
 - You cannot accurately reconstruct a signal that was sampled below its Nyquist frequency - you do not have the information
 - There is no point sampling at higher frequency - you do not gain extra information
- Signals that are not bandlimited cannot be accurately sampled and reconstructed
 - They would require an infinite sampling frequency


Sampling in Spatial Domain

- ❑ Sampling in the spatial domain is like multiplying by a spike function
- ❑ You take some ideal function and get data for a regular grid of points




Sampling in Frequency Domain

- Sampling in the frequency domain is like convolving with a spike function
 - Follows from the convolution theory: multiplication in spatial equals convolution in frequency
 - Spatial spike function in the frequency domain is also the spike function




Reconstruction (Frequency Domain)

- ❑ To reconstruct, we must restore the original spectrum
- ❑ That can be done by multiplying by a square pulse




Reconstruction (Spatial Domain)

- Multiplying by a square pulse in the frequency domain is the same as convolving with a sinc function in the spatial domain



Aliasing Due to Under-sampling


- If the sampling rate is too low, high frequencies get reconstructed as lower frequencies



- High frequencies from one copy get added to low frequencies from another

More Aliasing

- Poor reconstruction also results in aliasing
- Consider a signal reconstructed with a box filter in the spatial domain (square box pixels, which means using a sinc in the frequency domain):



Aliasing in Practice

- We have two types of aliasing:
 - Aliasing due to insufficient sampling frequency
 - Aliasing due to poor reconstruction
- You have some control over reconstruction
 - If resizing, for instance, use an approximation to the sinc function to reconstruct (instead of Bartlett, as we used last time)
 - Gaussian is closer to sinc than Bartlett
 - But note that sinc function goes on forever (*infinite support*), which is inefficient to evaluate
- You have some control over sampling if creating images using a computer
 - Remove all sharp edges (high frequencies) from the scene before drawing it
 - That is, blur character and line edges *before* drawing

Next Time

- ☐ Composition
- ☐ NPR
- ☐ 3D Graphics Toolkits