

Computer Graphics

Prof. Feng Liu

Fall 2016

<http://www.cs.pdx.edu/~fliu/courses/cs447/>

11/07/2016

Last time

☐ Hidden Surface Removal

Today

- ☐ Lighting and Shading
- ☐ Project 2
- ☐ Will publicize several times in the final week of classes when you can get your project graded
 - Demo your program to the instructor **in person**
 - ☐ Bring your own laptop or on a CS Windows Lab Machine
 - Latest time to grade
 - ☐ 5:00 pm, Friday, December 2, 2015
 - **No late submission!**

Where We Stand

- So far we know how to:
 - Transform between spaces
 - Draw polygons
 - Decide what's in front
 - Next
 - Deciding a pixel's intensity and color
-

Normal Vectors

- The intensity of a surface depends on its orientation with respect to the light and the viewer
 - The *surface normal vector* describes the orientation of the surface at a point
 - Mathematically: Vector that is perpendicular to the tangent plane of the surface
 - Just “the normal vector” or “the normal”
 - Will use \mathbf{n} or \mathbf{N} to denote
 - Normals are either supplied by the user or automatically computed
-

Transforming Normal Vectors

- Normal vectors are *directions*

Normal vectors are perpendicular to tangent vectors : $\mathbf{n} \bullet (\mathbf{x} - \mathbf{p}) = 0$

There is a matrix form of this : $\mathbf{n}^t (\mathbf{x} - \mathbf{p}) = 0$

Consider the equation with a transformed tangent : $\mathbf{n}^t \mathbf{T}^{-1} \mathbf{T} (\mathbf{x} - \mathbf{p}) = 0$

The right hand half is the transformed point.

The new transpose normal must be equal to : $\mathbf{n}^t \mathbf{T}^{-1}$

The new normal must then be : $(\mathbf{n}^t \mathbf{T}^{-1})^t = (\mathbf{T}^{-1})^t \mathbf{n}$


- To transform a normal, multiply it by the inverse transpose of the transformation matrix
 - Recall, rotation matrices are their own inverse transpose
 - Don't include the translation! Use $(n_x, n_y, n_z, 0)$ for homogeneous coordinates
-

Local Shading Models

- *Local shading models* provide a way to determine the intensity and color of a point on a surface
 - The models are local because they **don't consider other objects**
 - We use them because they are fast and simple to compute
 - They do not require knowledge of the entire scene, only the current piece of surface.
 - For the moment, assume:
 - We are applying these computations at a particular point on a surface
 - We have a normal vector for that point
-




Local Shading Models

- ❑ What they capture:
 - Direct illumination from light sources
 - *Diffuse* and *Specular* reflections
 - (Very) Approximate effects of global lighting
- ❑ What they don't do:
 - Shadows
 - Mirrors
 - Refraction
 - Lots of other stuff ...




“Standard” Lighting Model

- Consists of three terms linearly combined:
 - *Diffuse* component for the amount of incoming light from a point source reflected equally in all directions
 - *Specular* component for the amount of light from a point source reflected in a mirror-like fashion
 - *Ambient* term to approximate light arriving via other surfaces




Diffuse Illumination



$$k_d I_i (\mathbf{L} \cdot \mathbf{N})$$

- Incoming light, I_i from direction \mathbf{L} , is reflected equally in all directions
 - No dependence on viewing direction
 - Amount of light reflected depends on:
 - Angle of surface with respect to light source
 - Actually, determines how much light is collected by the surface, to then be reflected
 - Diffuse reflectance coefficient of the surface, k_d
 - Don't want to illuminate back side. Use $k_d I_i \max(\mathbf{L} \cdot \mathbf{N}, 0)$
-

Diffuse Example



Diffuse Lighting


Where is the light?

Which point is brightest
(how is the normal at the
brightest point related to
the light)?

Illustrating Shading Models

- Show the polar graph of the amount of light leaving for a given incoming direction:


Diffuse?



- Show the intensity of each point on a surface for a given light position or direction




Diffuse?



Specular Reflection


(Phong Reflectance Model)



$$k_s I_i (\mathbf{R} \bullet \mathbf{V})^p$$

- Incoming light is reflected primarily in the mirror direction, R
 - Perceived intensity depends on the relationship between the viewing direction, V , and the mirror direction
 - Bright spot is called a *specularity*
- Intensity controlled by:
 - The specular reflectance coefficient, k_s
 - The *Phong Exponent*, p , controls the apparent size of the specularity
 - Higher p , smaller highlight

Specular Example




Plus Specular Highlight

Illustrating Shading Models

- Show the polar graph of the amount of light leaving for a given incoming direction:


Specular?



- Show the intensity of each point on a surface for a given light position or direction




Specular?



Alternative Specular Reflection Model

$$\mathbf{H} = (\mathbf{L} + \mathbf{V}) / \|\mathbf{L} + \mathbf{V}\|$$

$$k_s I_i (\mathbf{H} \bullet \mathbf{N})^p$$



- Compute based on normal vector and “halfway” vector, H
-

Putting It Together

$$I = k_a I_a + I_i \left(k_d (\mathbf{L} \bullet \mathbf{N}) + k_s (\mathbf{H} \bullet \mathbf{N})^p \right)$$

- Global ambient intensity, I_a :
 - Gross approximation to light bouncing around of all other surfaces
 - Modulated by ambient reflectance k_a
 - Just sum all the terms
 - If there are multiple lights, sum contributions from each light
 - Several variations, and approximations ...
-

Color

$$I_r = k_{a,r} I_{a,r} + I_{i,r} \left(k_{d,r} (\mathbf{L} \bullet \mathbf{N}) + k_{s,r} (\mathbf{H} \bullet \mathbf{N})^n \right)$$

- Do everything for three colors, r, g and b
 - Note that some terms (the expensive ones) are constant
 - Using only three colors is an approximation, but few graphics practitioners realize it
 - k terms depend on wavelength, should compute for continuous spectrum
-

Approximations for Speed

- The viewer direction, V , and the light direction, L , depend on the surface position being considered, x
 - Distant light approximation:
 - Assume L is constant for all x
 - Good approximation if light is distant, such as sun
 - Distant viewer approximation
 - Assume V is constant for all x
 - Rarely good, but only affects specularities
-

Distant Light Approximation

- Distant light approximation:
 - Assume L is constant for all x
 - Good approximation if light is distant, such as sun
 - Generally called a *directional light source*
 - What aspects of surface appearance are affected by this approximation?
 - Diffuse?
 - Specular?
-

Distant Viewer Approximation

- Specularities require the viewing direction:
 - $V(x) = ||c-x||$
 - Slightly expensive to compute
 - Distant viewer approximation uses a global V
 - Independent of which point is being lit
 - Use the view plane normal vector
 - Error depends on the nature of the scene
 - Is the diffuse component affected?
-

Describing Surfaces

- The various parameters in the lighting equation describe the appearance of a surface
 - $(k_{d,r}, k_{d,g}, k_{d,b})$: The *diffuse color*, which most closely maps to what you would consider the “color” of a surface
 - Also called *diffuse reflectance coefficients*
 - $(k_{s,r}, k_{s,g}, k_{s,b})$: The specular color, which controls the color of specularities
 - Some systems do not let you specify this color separately
 - $(k_{a,r}, k_{a,g}, k_{a,b})$: The ambient color, which controls how the surface looks when not directly lit
 - Normally the same as the diffuse color
-

OpenGL Commands (1)

- ❑ `glMaterial{if}(face, parameter, value)`
 - Changes one of the coefficients for the front or back side of a face (or both sides)
 - ❑ `glLight{if}(light, property, value)`
 - Changes one of the properties of a light (intensities, positions, directions, etc)
 - There are 8 lights: `GL_LIGHT0`, `GL_LIGHT1`, ...
 - ❑ `glLightModel{if}(property, value)`
 - Changes one of the global light model properties (global ambient light, for instance)
 - ❑ `glEnable(GL_LIGHT0)` enables `GL_LIGHT0`
 - You must enable lights before they contribute to the image
 - You can enable and disable lights at any time
-

OpenGL Commands (2)

- ❑ `glEnable(GL_LIGHTING)` turns on lighting
 - You must enable lighting explicitly - it is off by default
 - ❑ Don't use specular intensity if you don't have to
 - It's expensive - turn it off by giving 0,0,0 as specular color of the lights
 - ❑ Don't forget normals
 - If you use scaling transformations, must enable `GL_NORMALIZE` to keep normal vectors of unit length
 - ❑ Many other things to control appearance
-

Light Sources

- Two aspects of light sources are important for a local shading model:
 - Where is the light coming from (the L vector)?
 - How much light is coming (the I values)?
 - Various light source types give different answers to the above questions:
 - *Point light source*: Light from a specific point
 - *Directional*: Light from a specific direction
 - *Spotlight*: Light from a specific point with intensity that depends on the direction
 - *Area light*: Light from a continuum of points (later in the course)
-

Point and Directional Sources

□ Point light: $L(x) = \frac{p_{light} - x}{\|p_{light} - x\|}$

- The L vector depends on where the surface point is located
- Must be normalized - slightly expensive
- To specify an OpenGL light at 1,1,1:


```
GLfloat light_position[] = { 1.0, 1.0, 1.0, 1.0 };  
glLightfv(GL_LIGHT0, GL_POSITION, light_position);
```

□ Directional light: $L(x) = L_{light}$

- The L vector does not change over points in the world
- OpenGL light traveling in direction 1,1,1 (L is in opposite direction):

```
GLfloat light_position[] = { 1.0, 1.0, 1.0, 0.0 };  
glLightfv(GL_LIGHT0, GL_POSITION, light_position);
```

Spotlights



□ Point source, but intensity depends on L :

- Requires a position: the location of the source

```
glLightfv(GL_LIGHT0, GL_POSITION, light_posn);
```

- Requires a direction: the center axis of the light

```
glLightfv(GL_LIGHT0, GL_SPOT_DIRECTION, light_dir);
```

- Requires a cut-off: how broad the beam is

```
glLightfv(GL_LIGHT0, GL_SPOT_CUTOFF, 45.0);
```

- Requires an exponent: how the light tapers off at the edges of the cone

- Intensity scaled by $(L \cdot D)^n$

```
glLightfv(GL_LIGHT0, GL_SPOT_EXPONENT, 1.0);
```

Shading so Far

- So far, we have discussed illuminating a single point

$$I = k_a I_a + I_i \left(k_d (\mathbf{L} \bullet \mathbf{N}) + k_s (\mathbf{H} \bullet \mathbf{N})^p \right)$$


- We have assumed that we know:
 - The point
 - The surface normal
 - The viewer location (or direction)
 - The light location (or direction)
 - But commonly, normal vectors are only given at the vertices
 - It is also expensive to compute lighting for every point
-

Shading Interpolation

- Take information specified or computed at the vertices, and somehow propagate it across the polygon (triangle)
 - Several options:
 - Flat shading
 - Gouraud interpolation
 - Phong interpolation
-


Flat shading

- ❑ Compute shading at a representative point and apply to whole polygon
 - OpenGL uses one of the vertices
- ❑ Advantages:
 - Fast - one shading computation per polygon
- ❑ Disadvantages:
 - Inaccurate
 - What are the artifacts?




Gouraud Shading

- Shade each *vertex* with it's own location and normal
- *Linearly interpolate* the color across the face
- Advantages:
 - Fast: incremental calculations when rasterizing
 - Much smoother - use same normal every time a vertex is used for a face
- Disadvantages:
 - What are the artifacts?
 - Is it accurate?




Phong Interpolation


- ❑ Interpolate normals across faces
- ❑ Shade each pixel individually
- ❑ Advantages:
 - High quality, narrow specularities
- ❑ Disadvantages:
 - Expensive
 - Still an approximation for most surfaces
- ❑ Not to be confused with Phong's specular model



Gouraud



Phong



Shading and OpenGL

- ❑ OpenGL defines two particular shading models
 - Controls how colors are assigned to pixels
 - `glShadeModel(GL_SMOOTH)` interpolates between the colors at the vertices (the default, Gouraud shading)
 - `glShadeModel(GL_FLAT)` uses a constant color across the polygon
 - ❑ Phong shading requires a significantly greater programming effort - beyond the scope of this class
 - Also requires *fragment shaders* on *programmable graphics hardware*
-

Next Time

☐ Texture mapping